

CSET419 – Introduction to Generative AI

Lab – 4: Text Generation Models

Submitted By: Deepanshu
Date: 2026-02-04

Objective

The objective of this lab is to design and implement a simple text generation model that can learn patterns from a given text corpus and generate new, meaningful text sequences.

Learning Outcomes

- Understand the basics of text generation
- Preprocess textual data for neural networks
- Implement a sequence-based neural network model (RNN/LSTM/GRU)
- Implement a Transformer-based model
- Generate new text using a trained model

Component-I: RNN / LSTM Based Text Generation

This section demonstrates text generation using Recurrent Neural Networks (RNN), Long Short-Term Memory (LSTM), and Gated Recurrent Units (GRU).

1. Dataset & Preprocessing

A custom text corpus was used. Tokenization implies character-level mapping.

Dataset Snippet:

```
artificial intelligence is transforming modern society.  
it is used in healthcare finance education and transportation.  
machine learning allows systems to improve automatically with experience.  
data plays a critical role in training intelligent systems.  
large datasets help models learn complex patterns.  
...
```

2. Simple RNN Implementation

Implementation of a vanilla RNN using `nn.RNN`.

Code (rnn_gen.py):

```
import torch  
import torch.nn as nn  
import torch.optim as optim  
import os  
import utils  
  
MODEL_FILE = 'rnn_model.pth'  
SEQ_LENGTH = 40  
HIDDEN_SIZE = 128  
NUM_LAYERS = 1  
  
class TextGeneratorRNN(nn.Module):  
    def __init__(self, vocab_size, hidden_size, num_layers):  
        super(TextGeneratorRNN, self).__init__()  
        self.embedding = nn.Embedding(vocab_size, hidden_size)  
        self.rnn = nn.RNN(hidden_size, hidden_size, num_layers, batch_first=True)  
        self.fc = nn.Linear(hidden_size, vocab_size)  
  
    def forward(self, x, hidden):  
        embed = self.embedding(x)  
        out, hidden = self.rnn(embed, hidden)  
        out = self.fc(out[:, -1, :])  
        return out, hidden  
  
    def train_model():  
        text, dataX, dataY, char_to_int, _, vocab_size = utils.load_data(SEQ_LENGTH)  
        if text is None: return  
  
        X = torch.tensor(dataX, dtype=torch.long)  
        y = torch.tensor(dataY, dtype=torch.long)  
        dataset = torch.utils.data.TensorDataset(X, y)  
        dataloader = torch.utils.data.DataLoader(dataset, batch_size=64, shuffle=True)  
  
        device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')  
        if torch.backends.mps.is_available(): device = torch.device('mps')  
  
        model = TextGeneratorRNN(vocab_size, HIDDEN_SIZE, NUM_LAYERS).to(device)  
        criterion = nn.CrossEntropyLoss()  
        optimizer = optim.Adam(model.parameters(), lr=0.005)  
  
        print("Starting Training (Vanilla RNN)...")  
        for epoch in range(100):  
            model.train()  
            total_loss = 0  
            for batch_x, batch_y in dataloader:
```

```

batch_x, batch_y = batch_x.to(device), batch_y.to(device)
optimizer.zero_grad()
hidden = None
output, _ = model(batch_x, hidden)
loss = criterion(output, batch_y)
loss.backward()
optimizer.step()
total_loss += loss.item()

if (epoch + 1) % 10 == 0:
    print(f"Epoch {epoch+1}/100, Loss: {total_loss/len(dataloader):.4f}")

torch.save(model.state_dict(), MODEL_FILE)
print(f"RNN model saved to {MODEL_FILE}")

def generate_text_model(start_str, length=200):
    if not os.path.exists(MODEL_FILE):
        print("Model not found. Please train first.")
        return ""

    _, _, _, char_to_int, int_to_char, vocab_size = utils.load_data(SEQ_LENGTH)

    device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
    if torch.backends.mps.is_available(): device = torch.device('mps')

    model = TextGeneratorRNN(vocab_size, HIDDEN_SIZE, NUM_LAYERS).to(device)
    model.load_state_dict(torch.load(MODEL_FILE, map_location=device))
    model.eval()

    current_string = start_str.lower()
    generated_text = start_str

    with torch.no_grad():
        for _ in range(length):
            if len(current_string) >= SEQ_LENGTH: input_seq = current_string[-SEQ_LENGTH:]
            else: input_seq = current_string

            input_indices = [char_to_int.get(c, 0) for c in input_seq]
            input_tensor = torch.tensor(input_indices, dtype=torch.long).unsqueeze(0).to(device)

            output, _ = model(input_tensor, None)
            probs = torch.softmax(output, dim=1)
            next_char_idx = torch.multinomial(probs, 1).item()

            next_char = int_to_char[next_char_idx]
            generated_text += next_char
            current_string += next_char

    return generated_text

if __name__ == "__main__":
    train_model()

```

Generated Output (RNN):

artificial intelligent systems.
large datasets help models are important concerns.
researchers continue to improve ai systems should be meaning.
automatically with experiences.

ethical consid on languages trai

3. LSTM & GRU Implementation

LSTM and GRU architectures address the vanishing gradient problem in simple RNNs.

Code (lstm_gru_gen.py):

```
import torch
import torch.nn as nn
import torch.optim as optim
import os
import utils

SEQ_LENGTH = 40
HIDDEN_SIZE = 256
NUM_LAYERS = 2

class TextGeneratorLSTM(nn.Module):
    def __init__(self, vocab_size, hidden_size, num_layers):
        super(TextGeneratorLSTM, self).__init__()
        self.embedding = nn.Embedding(vocab_size, hidden_size)
        self.lstm = nn.LSTM(hidden_size, hidden_size, num_layers, batch_first=True)
        self.fc = nn.Linear(hidden_size, vocab_size)

    def forward(self, x, hidden):
        embed = self.embedding(x)
        out, hidden = self.lstm(embed, hidden)
        out = self.fc(out[:, -1, :])
        return out, hidden

class TextGeneratorGRU(nn.Module):
    def __init__(self, vocab_size, hidden_size, num_layers):
        super(TextGeneratorGRU, self).__init__()
        self.embedding = nn.Embedding(vocab_size, hidden_size)
        self.gru = nn.GRU(hidden_size, hidden_size, num_layers, batch_first=True)
        self.fc = nn.Linear(hidden_size, vocab_size)

    def forward(self, x, hidden):
        embed = self.embedding(x)
        out, hidden = self.gru(embed, hidden)
        out = self.fc(out[:, -1, :])
        return out, hidden

def train_model(model_type='LSTM'):
    text, dataX, dataY, _, _, vocab_size = utils.load_data(SEQ_LENGTH)
    if text is None: return

    X = torch.tensor(dataX, dtype=torch.long)
    y = torch.tensor(dataY, dtype=torch.long)
    dataset = torch.utils.data.TensorDataset(X, y)
    dataloader = torch.utils.data.DataLoader(dataset, batch_size=64, shuffle=True)

    device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
    if torch.backends.mps.is_available(): device = torch.device('mps')

    if model_type == 'LSTM':
        model = TextGeneratorLSTM(vocab_size, HIDDEN_SIZE, NUM_LAYERS).to(device)
        model_file = 'lstm_model.pth'
    else:
        model = TextGeneratorGRU(vocab_size, HIDDEN_SIZE, NUM_LAYERS).to(device)
        model_file = 'gru_model.pth'

    criterion = nn.CrossEntropyLoss()
    optimizer = optim.Adam(model.parameters(), lr=0.002)

    print(f"Starting Training ({model_type})...")
    for epoch in range(100):
        model.train()
        total_loss = 0
        for batch_x, batch_y in dataloader:
```

```

batch_x, batch_y = batch_x.to(device), batch_y.to(device)
optimizer.zero_grad()
hidden = None
output, _ = model(batch_x, hidden)
loss = criterion(output, batch_y)
loss.backward()
optimizer.step()
total_loss += loss.item()

if (epoch + 1) % 10 == 0:
    print(f"Epoch {epoch+1}/100, Loss: {total_loss/len(dataloader):.4f}")

torch.save(model.state_dict(), model_file)
print(f"{model_type} model saved to {model_file}")

def generate_text_model(start_str, model_type='LSTM', length=200):
    model_file = 'lstm_model.pth' if model_type == 'LSTM' else 'gru_model.pth'

    if not os.path.exists(model_file):
        print(f"Model {model_file} not found. Please train first.")
        return ""

    _, _, _, char_to_int, int_to_char, vocab_size = utils.load_data(SEQ_LENGTH)

    device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
    if torch.backends.mps.is_available(): device = torch.device('mps')

    if model_type == 'LSTM':
        model = TextGeneratorLSTM(vocab_size, HIDDEN_SIZE, NUM_LAYERS).to(device)
    else:
        model = TextGeneratorGRU(vocab_size, HIDDEN_SIZE, NUM_LAYERS).to(device)

    model.load_state_dict(torch.load(model_file, map_location=device))
    model.eval()

    current_string = start_str.lower()
    generated_text = start_str

    with torch.no_grad():
        for _ in range(length):
            if len(current_string) >= SEQ_LENGTH: input_seq = current_string[-SEQ_LENGTH:]
            else: input_seq = current_string

            input_indices = [char_to_int.get(c, 0) for c in input_seq]
            input_tensor = torch.tensor(input_indices, dtype=torch.long).unsqueeze(0).to(device)

            output, _ = model(input_tensor, None)
            probs = torch.softmax(output, dim=1)
            next_char_idx = torch.multinomial(probs, 1).item()

            next_char = int_to_char[next_char_idx]
            generated_text += next_char
            current_string += next_char

    return generated_text

if __name__ == "__main__":
    train_model('LSTM')

```

Generated Output (LSTM):

artificial intelligence.
 fairness transparency and accountability must be ensured.
 ai systems should be designed responsibly.
 data privacy and security are major concerns.
 researchers continue to improve ai saf

Generated Output (GRU):

artificial intelligence.

fairness transparency and accountability must be ensured.
ai systems should be designed responsibly.
data privacy and security are major concerns.
researchers continue to improve ai saf

Component-II: Transformer Based Text Generation

Transformer models rely on self-attention mechanisms to process sequential data in parallel, offering significant performance improvements over RNNs for many tasks.

Implementation Details

Uses `nn.TransformerEncoder` with Positional Encoding to retain sequence order information.

Code (transformer_gen.py):

```
import torch
import torch.nn as nn
import torch.optim as optim
import math
import os
import utils

MODEL_FILE = 'transformer_model.pth'
SEQ_LENGTH = 40
D_MODEL = 128
NHEAD = 4
NUM_LAYERS = 2

class PositionalEncoding(nn.Module):
    def __init__(self, d_model, max_len=5000):
        super(PositionalEncoding, self).__init__()
        pe = torch.zeros(max_len, d_model)
        position = torch.arange(0, max_len, dtype=torch.float).unsqueeze(1)
        div_term = torch.exp(torch.arange(0, d_model, 2).float() * (-math.log(10000.0) / d_model))
        pe[:, 0::2] = torch.sin(position * div_term)
        pe[:, 1::2] = torch.cos(position * div_term)
        pe = pe.unsqueeze(0)
        self.register_buffer('pe', pe)

    def forward(self, x):
        return x + self.pe[:, :x.size(1), :]

class TextGeneratorTransformer(nn.Module):
    def __init__(self, vocab_size, d_model, nhead, num_layers):
        super(TextGeneratorTransformer, self).__init__()
        self.embedding = nn.Embedding(vocab_size, d_model)
        self.pos_encoder = PositionalEncoding(d_model)
        encoder_layers = nn.TransformerEncoderLayer(d_model, nhead, batch_first=True)
        self.transformer_encoder = nn.TransformerEncoder(encoder_layers, num_layers)
        self.fc = nn.Linear(d_model, vocab_size)
        self.d_model = d_model

    def forward(self, x):
        embed = self.embedding(x) * math.sqrt(self.d_model)
        src = self.pos_encoder(embed)
        output = self.transformer_encoder(src)
        out = self.fc(output[:, -1, :])
        return out

    def train_model():
        text, dataX, dataY, _, _, vocab_size = utils.load_data(SEQ_LENGTH)
        if text is None: return

        X = torch.tensor(dataX, dtype=torch.long)
        y = torch.tensor(dataY, dtype=torch.long)
        dataset = torch.utils.data.TensorDataset(X, y)
        dataloader = torch.utils.data.DataLoader(dataset, batch_size=64, shuffle=True)

        device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
        if torch.backends.mps.is_available(): device = torch.device('mps')

        model = TextGeneratorTransformer(vocab_size, D_MODEL, NHEAD, NUM_LAYERS).to(device)
```

```

criterion = nn.CrossEntropyLoss()
optimizer = optim.Adam(model.parameters(), lr=0.0005)

print("Starting Training (Transformer)...")
for epoch in range(100):
    model.train()
    total_loss = 0
    for batch_x, batch_y in dataloader:
        batch_x, batch_y = batch_x.to(device), batch_y.to(device)
        optimizer.zero_grad()
        output = model(batch_x)
        loss = criterion(output, batch_y)
        loss.backward()
        optimizer.step()
        total_loss += loss.item()

    if (epoch + 1) % 10 == 0:
        print(f"Epoch {epoch+1}/100, Loss: {total_loss/len(dataloader):.4f}")

torch.save(model.state_dict(), MODEL_FILE)
print(f"Transformer model saved to {MODEL_FILE}")

def generate_text_model(start_str, length=200):
    if not os.path.exists(MODEL_FILE):
        print("Model not found. Please train first.")
        return ""

    _, _, char_to_int, int_to_char, vocab_size = utils.load_data(SEQ_LENGTH)

    device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
    if torch.backends.mps.is_available(): device = torch.device('mps')

    model = TextGeneratorTransformer(vocab_size, D_MODEL, NHEAD, NUM_LAYERS).to(device)
    model.load_state_dict(torch.load(MODEL_FILE, map_location=device))
    model.eval()

    current_string = start_str.lower()
    generated_text = start_str

    with torch.no_grad():
        for _ in range(length):
            if len(current_string) >= SEQ_LENGTH: input_seq = current_string[-SEQ_LENGTH:]
            else: input_seq = current_string

            input_indices = [char_to_int.get(c, 0) for c in input_seq]
            input_tensor = torch.tensor(input_indices, dtype=torch.long).unsqueeze(0).to(device)

            output = model(input_tensor)

            probs = torch.softmax(output, dim=1)
            next_char_idx = torch.multinomial(probs, 1).item()

            next_char = int_to_char[next_char_idx]
            generated_text += next_char
            current_string += next_char

    return generated_text

if __name__ == "__main__":
    train_model()

```

Generated Output (Transformer):

artificial trrnt sp.
 t utfove an ange agen as agelatines cy leleritemodetems parntes.
 isthatanthint arngrexlli mperstalalarnsterncct mpattate matatroncles.
 tentowscy sso sto she to tonitysy mon bon alon .
 a qua

Conclusion

In this lab, we successfully implemented and compared N-gram, RNN, LSTM, GRU, and Transformer models for character-level text generation. While simple RNNs can generate text, LSTM and GRU models show better capability in capturing longer dependencies. Transformer models, although more complex, provide a powerful architecture for sequence tasks.